

# Introduction to Graph Databases

Swaminathan Muthuveerappan

# Agenda

- Why Graph Databases?
- Neo4j's Cypher Query Language (CQL)
- Optimization: Indexing
- Optimization: Procedures
- Optimization: Streaming
- Optimization: Queryable Relationships

# What is a Database?

- Organize data for easy access
- Fast searching of documents

# Types of Databases

- SQL Databases (MySQL, PostgreSQL, Google CloudSQL etc...)
- NoSQL Databases (MongoDb, Google Datastore etc...)
- Graph Databases (Neo4j, ArangoDB)

**Let's design a database for an  
ecommerce application (promotions  
domain)**

# Specifications

1. Business should be able to create Promotions on their inventory
2. Business should be able to know the list of applicable SKUs/Variants for a given Promotion

*SKU - Stock Keeping Unit*

# Inventory Specification

- Category
  - Mobile Phones, Electronics etc...
- Product
  - iPhone, LG OLED TV etc...
- Variant
  - iPhone 128GB Space Grey, LG Smart OLED 4K 65B9PUA etc...

*(image)*

**Will SQL Databases work well?**



# **How about NoSQL Databases?**

**Finally...**

# Introduction to Graph Databases :)

# Why would Graph Databases work?

- Closely represent whiteboard models
- Relationships are first-class citizens
- Cheap Traversals

*(image)*

# The Cypher Query Language

# **Let's create a simple graph using Cypher queries**

# CREATE Clause

```
CREATE (appleBrand:Brand{id:'apple_123', name:'Apple'}) RETURN appleBrand;
```

# Relationships

```
MATCH (appleBrand:Brand{id:'apple_123'})  
MATCH (iPhoneVariant:Variant{id:'variant_123'})  
CREATE (appleBrand)-[:has_variant]->(iPhoneVariant);
```



# Querying

Get all the variants of a given promotion

*(image)*

# Optimizations

# **1. Indexing to improve search performance**

## **2. Extending Neo4j with User-Defined procedures**

# **3. Streaming for large data**

## **4. Improving performance with proper Relationship names**

**Thank you for your time**

:)



# Questions?